

ASSIGNMENT 1 - OPERATING SYSTEMS LAB CS344

Group C6 :

1. Vani Krishna Barla 200101101
2. Mahek Vora 200101062
3. Siddharth Bansal 200101093

Part 1: Kernel threads

POSIX thread libraries(thread APIs for C and C++) allow us to spawn a new current process flow. This library is used in software for faster execution.

Thread creation requires lesser overhead than forking as thread creation does not initialize new system virtual memory space and environment for the process. Further, all threads share the same address space

System calls `sys_thread_create`, `sys_thread_join`, `sys_thread_exit` are created in `syscall.c` that act as an interface for the user to implement `thread_create`, `thread_join` and `thread_exit` functions in `proc.c` respectively by taking arguments and passing it to `proc.c` functions.

```
int sys_thread_create(void) {
    int func_add;
    int arg;
    int stack_add;

    if (argint(0, &func_add) < 0)
        return -1;
    if (argint(1, &arg) < 0)
        return -1;
    if (argint(2, &stack_add) < 0)
        return -1;

    return thread_create((void *)func_add, (void *)arg, (void *)stack_add);
}

int sys_thread_join(void) {
    return thread_join();
}

int sys_thread_exit(void) {
    return thread_exit();
}
```

l.thread_create()

```
struct proc {
    int is_thread;
    void *stack;
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];         // Process name (debugging)
};
```

proc.h

- Structurally, process and thread are similar apart from the fact that threads share the same address space
- So for threads we can modify the structure “**proc**” in “**proc.h**”
 - Add a flag variable **int “is_thread”** to indicate if the structure is a thread (is_thread=1) or a process (is_thread=0)
 - Add a “**void* stack**” variable that points to the stack memory of the process.

```
int thread_create(void(*fcn)(void*), void *arg, void*stack) {
    int i, pid;
    struct proc *np;
    struct proc *curproc = myproc();
    // int *myarg;
    // int *myret;
    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }

    np->sz = curproc->sz;
    np->parent = curproc;
```

```

np->pgdir = curproc->pgdir;
*np->tf = *curproc->tf;
np->is_thread = 1;
np->tf->eax = 0;
np->tf->eip = (int) fcn; // instruction pointer
np->stack = stack;
np->tf->esp = (uint)stack + 4092; // stack pointer
*((uint *) (np->tf->esp)) = (uint)arg;
*((uint *) (np->tf->esp - 4)) = 0xFFFFFFFF;
np->tf->esp -= 4;
for(i = 0; i < NOFILE; i++)
    if(curproc->ofile[i])
        np->ofile[i] = filedup(curproc->ofile[i]);
np->cwd = idup(curproc->cwd);
safestrcpy(np->name, curproc->name, sizeof(curproc->name));
pid = np->pid;
acquire(&ptable.lock);
np->state = RUNNABLE;
release(&ptable.lock);
return pid;
}

```

proc.c

A new function called *thread_create* is added here. It functions similarly to fork (which creates a child process).

Working:

- **allocproc()** looks into the process table for “unused proc” and if found, the state is changed to “EMBRYO”, if not found, 0 is returned by allocproc() and thread_create() returns -1.
- **sz**: size of process/thread memory = same as parent process as they share the address space
- **parent**: is set to current process
- The thread will have the same process state as the current process. So “pgdir” (i.e. mapping in physical and virtual addresses) will be the same
- **tf**: (trap frame, it holds the userspace state) will be the same as that of the current process
- **is_thread**: set to true
- **eip**: The instruction pointer is set to the function which is taken as an input parameter of the function as required
- **stack**: the current process’s stack variable is initialized using the pointer taken as an input parameter of the function. As the esp points to the top of the stack, and 2

locations(the size of int*) are reserved, taking the page size as 4096, the esp should point at **starting_location + 4096 - 2*(sizeof(int*))**
that is, **(uint)stack + 4092**

- **esp:** (stack pointer) 2 locations in the stack are reserved for arguments and the fake address respectively
- **eax register:** is cleared
- **Open files:** of the current process are duplicated for the newly created thread
- **Current working directory** of the thread is set to the cwd of the parent process, thus bumping up the reference count of parent's cwd by calling `idup()`.
- **Process name:** will also remain same for the new thread.
- The process state is set to `RUNNABLE` and this is done in the critical section as we do not want any context switches during the execution of this set of instructions
- It returns the new thread's pid.

II. `thread_join()`

```
int thread_join(void) {
    struct proc *p;
    int havekids, pid;
    struct proc *curproc = myproc();
    acquire(&ptable.lock);
    for (;;) {
        havekids = 0;
        for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
            if (p->parent != curproc || p->is_thread != 1)
                continue;
            havekids = 1;
            if (p->state == ZOMBIE) {
                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                p->state = UNUSED;
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                release(&ptable.lock);
                return pid;
            }
        }
    }

    if (!havekids || curproc->killed) {
```

```

    release(&ptable.lock);
    return -1;
}
sleep(curproc, &ptable.lock);
}
}

```

proc.c

- *thread_join* function waits for a spawned thread to exit and return its pid. If there are no spawned threads, it returns -1.
- It first checks if any process is its spawned thread (i.e. it is a child of the current process and is a thread).
- If the process state is ZOMBIE, perform cleanup and return. If a thread exists but is not in ZOMBIE state, the process goes to sleep state.
- The function *thread_join()* code is referenced from the function *wait()*. The code for which is mostly similar to *wait* except for two things.
 - It checks `p->is_thread == 1` along with making sure the process *p* is the child of the current process, since we want to distinguish the thread process.
 - It doesn't free up the page directory of the thread as it is shared across all the threads of the process as well as the parent process.

III. thread_exit()

```

int thread_exit(void) {
    struct proc *p;
    struct proc *curproc = myproc();
    int fd;

    if(curproc == initproc)
        panic("init exiting");

    for (fd = 0; fd < NOFILE; fd++) {
        if (curproc -> ofile[fd]) {
            fileclose(curproc -> ofile[fd]);
            curproc -> ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc -> cwd);
    end_op();
    curproc -> cwd = 0;
}

```

```

acquire(&ptable.lock);

// Parent might be sleeping in wait().
wakeup1(curproc->parent);

// Pass abandoned children to init.
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->parent == curproc){
        p->parent = initproc;
        if(p->state == ZOMBIE)
            wakeup1(initproc);
    }
}

// Jump into the scheduler, never to return.
curproc->state = ZOMBIE;
sched();
panic("zombie exit");
}

```

proc.c

A new process called *thread_exit* is added, it is structurally similar to *exit*

Working:

- The process that is exiting must not be the *initproc*
- Close all the open files for the current thread and setting "ofile" to 0 indicates we are no longer accessing the file
- Call *wakeup1* on the parent process that might be sleeping
- Set the current process state to "ZOMBIE" and call *sched*, at this point it holds the *ptable* lock

Part 2: Synchronization

Synchronization error can be avoided by surrounding the atomic section of the code with locks. Here the atomic section in **do_work()** function:

```
old = total_balance;
delay(100000);
total_balance = old + 1;
```

which is accessing a variable and updating its value, a delay is added to increase chances of context-switching between threads accessing the same variable and updating it unsynchronized.

I. Spinlocks

Creating a struct in C to define the structure of lock which contains two variables as given below, *locked* signifies the state of the lock initialized and *name* is the name of the lock added for debugging purposes to distinguish between locks.

```
struct spinlock {
    uint locked;        // Is the lock held?
    char *name;         // Name of lock.
};
```

The *init* function initializes the lock's value of *locked* (*lk->locked*) to 0 signifying this lock is not yet locked/ is available for use.

```
void
thread_spin_init(struct spinlock *lk, char *name)
{
    lk->name = name;
    lk->locked = 0;
}
```

The *thread_spin_lock()* function is used to acquire the lock only if the lock is available (i.e. *lk->locked* value is 0) and to show it is acquired and is no longer available, the function sets the value of *lk->locked* = 1.

This is achieved by calling the **atomic function xchg(&lk->locked, 1)** which assigns the value of *locked* to 1 and returns the *original value* of *locked*. If the returned value is *not equal* to 0 signifies that it was already 1 before the function call, i.e. the lock has been acquired by some other thread.

Thus for this thread, the function *xchg* is called in a while loop until the lock is available, by which the original value of *lk->locked* returned would be 0 (*busy waiting*). Once the thread comes out of the while loop, it would have acquired the lock (the current value of *lk->locked* set to 1).

```
// Acquire the lock.
void
thread_spin_lock(struct spinlock *lk)
{
    while(xchg(&lk->locked, 1) != 0)
        ;
    __sync_synchronize();
}
```

While releasing the acquired lock, the value of `lk->locked` is set back to 0. This code segment uses `asm` (*inline assembly code* where `volatile` prevents an `asm` instruction from being deleted) instead of C assignment to make sure the value assignment operation is *atomic*.

```
// Release the lock.
void
thread_spin_unlock(struct spinlock *lk)
{
    __sync_synchronize();
    asm volatile("movl $0, %0" : "+m" (lk->locked) : );
}
```

The function `__sync_synchronize` is an atomic builtin used to force a full memory barrier. A full memory barrier prevents memory operations from being moved across it during compiler code optimization, neither forward nor backward.

It also stops the cpu from scheduling reads and writes, issued before the instruction, after the instruction and those issued after the instruction from being scheduled before it. **`__sync_synchronize`** is used here to prevent the memory operations in the critical section from being scheduled before acquiring the lock/ after releasing the lock

II. Mutex Locks

Structure for mutex lock is quite similar to the *spinlock* where we use an integer locked to signify the state of the mutex lock (0 if not acquired by any process and 1 if acquired by a process/thread)

```
struct thread_mutex {
    uint locked;        // Is the lock held?
};
```

Initializing the mutex lock by setting the locked value to 0, signifies that lock is available to be acquired.

```
void thread_mutex_init(struct thread_mutex *m) {
```



```

    m->locked = 0;
}

```

Acquiring the lock if it is available (that is, `xchg(&m->locked, 1)` returns 0) if not then calls `sleep(1)`.

The `sleep()` function shall cause the calling thread to be suspended from execution until the number of real time seconds specified by the argument `seconds` has elapsed.

This avoids the *busy waiting state* and allows the CPU to suspend the current thread and thus allowing the processor to be used for executing some other process/thread. This also avoids the scenario where all threads of the process start to spin endlessly waiting for lock to be released.

```

void thread_mutex_lock(struct thread_mutex *m)
{
    while(xchg(&m->locked, 1) != 0)
        sleep(1);
    __sync_synchronize();
}

```

The mutex lock is unlocked similarly to the spinlock by setting the value of `m->locked` to 0 using an inline assembly command to make sure it's an atomic operation.

The utilization of `__sync_synchronize()` function is similar to that mentioned above.

```

void thread_mutex_unlock(struct thread_mutex *m)
{
    //  unlock(m);
    __sync_synchronize();
    asm volatile("movl $0, %0" : "+m" (m->locked) : );
}

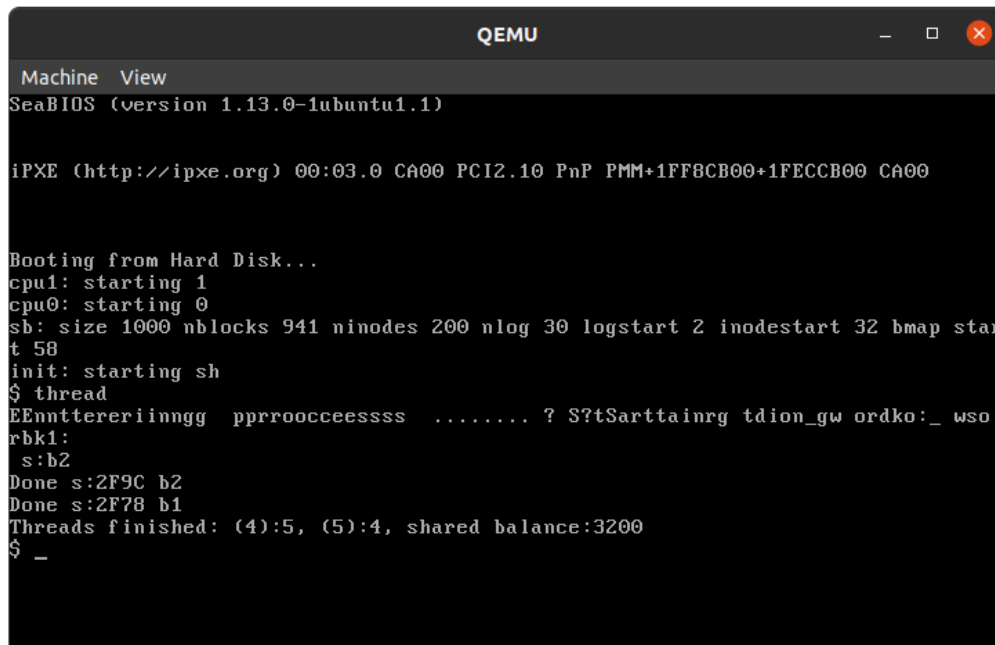
```

Note: `pushcli()` and `popcli()` functions used in spinlock implementation of xv6 (`spinlock.c/h`) from which the above code of locks is referenced from is not used because they are for disabling and enabling interrupts at OS / kernel level.

Whereas the spinlock and mutex lock implemented by us is used at user level thus these locks won't be used by interrupts and scenarios of deadlock can't arise.

Implementing locks and testing with `thread.c`

- I. **With no locks implemented** - since there is no synchronization, the balance value is being overwritten, thus the final total balance is not the correct value.

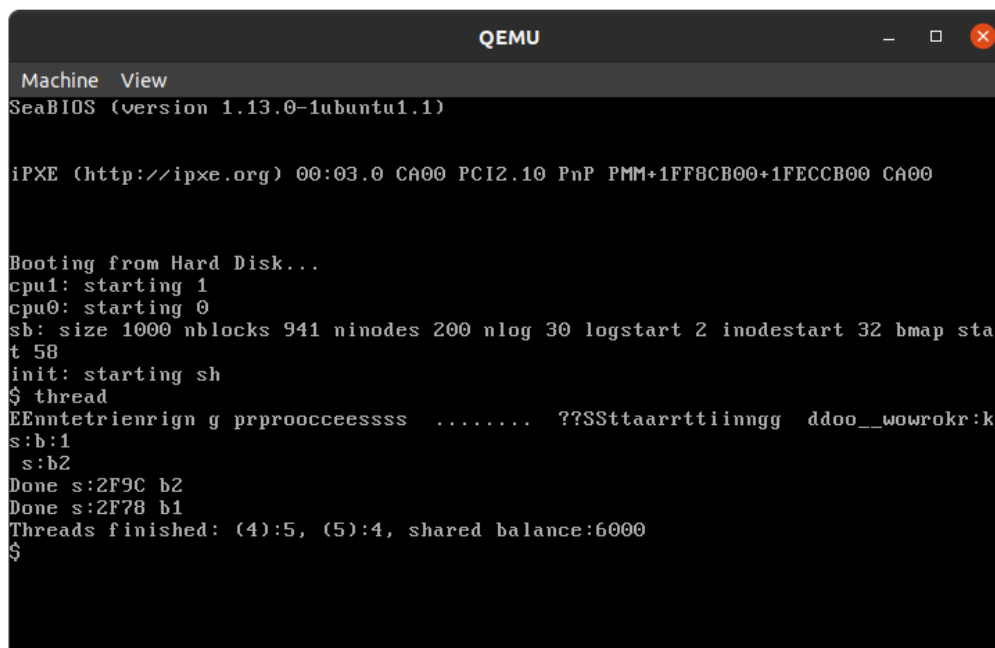


```
Machine View
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CB00+1FECCB00 CA00

Booting from Hard Disk...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ thread
EEEnnttereriinnngg pprroocceessss ..... ? S?tSarttainrg tdion_gw ordko:_ wso:
rbk1:
s:b2
Done s:2F9C b2
Done s:2F78 b1
Threads finished: (4):5, (5):4, shared balance:3200
$ _
```

- II. **Spinlock implemented with 2 processors** - since locks are implemented, the two threads do not enter the critical section at the same time and hence do not overwrite the balance while the other process is trying to increment the balance and thus the final total balance is the correct value($6000 = 3200 + 2800$).



```
Machine View
SeaBIOS (version 1.13.0-1ubuntu1.1)

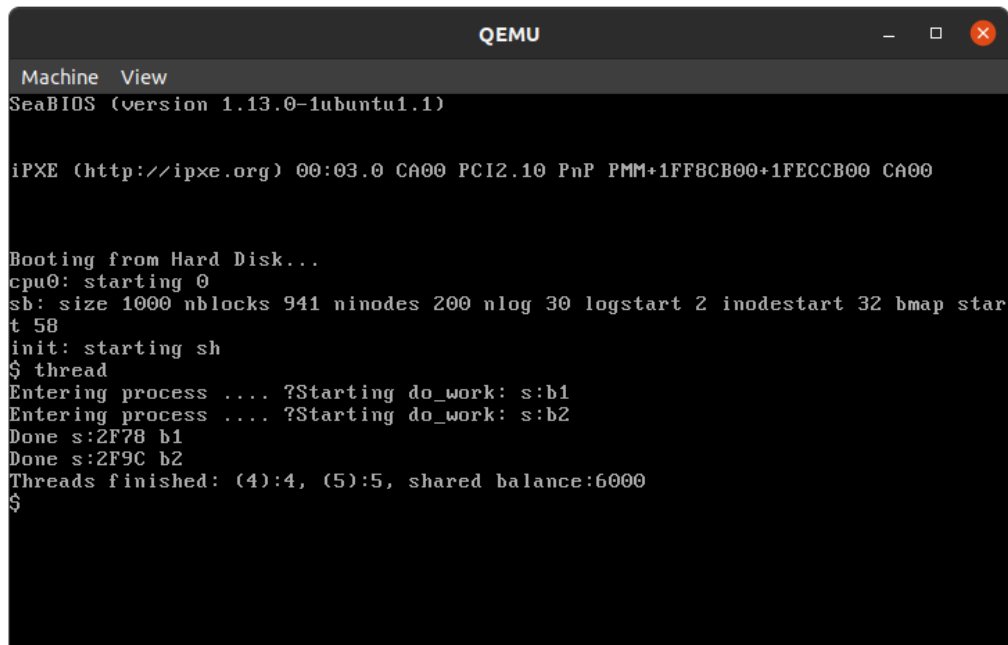
iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CB00+1FECCB00 CA00

Booting from Hard Disk...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ thread
EEEnnttetrienrign g pprroocceessss ..... ??SSttaarrttiinnngg ddoo__wowrokr:k
s:b:1
s:b2
Done s:2F9C b2
Done s:2F78 b1
Threads finished: (4):5, (5):4, shared balance:6000
$
```

The output balance remains the same in the implementation of both spin lock and mutex lock as the critical section of the `do_work` function is not significantly time consuming. As

this is the only process running on the os, the processors are not under heavy load, the efficiency of mutex locks over spin locks is not noticeable in this test case.

III. Mutex lock implemented with 1 processor

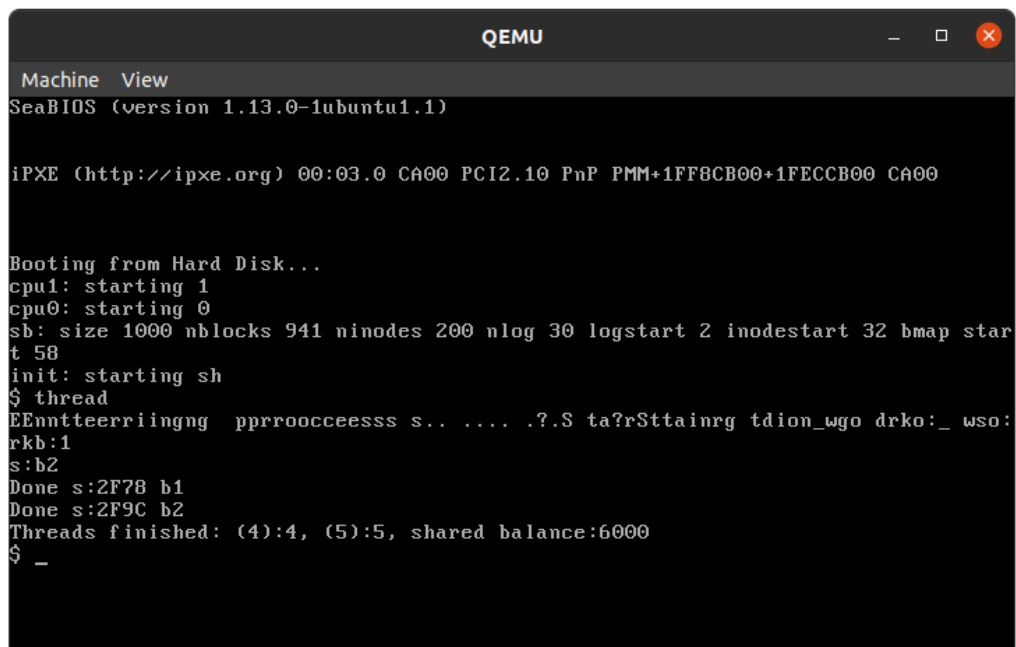


```
QEMU
Machine View
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CB00+1FECCB00 CA00

Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ thread
Entering process .... ?Starting do_work: s:b1
Entering process .... ?Starting do_work: s:b2
Done s:2F78 b1
Done s:2F9C b2
Threads finished: (4):4, (5):5, shared balance:6000
$
```

IV. Mutex lock with 2 processors



```
QEMU
Machine View
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CB00+1FECCB00 CA00

Booting from Hard Disk...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ thread
EEenntteerriingng pprroocceesss s.. .... ?.S ta?rSttainrg tdion_wgo drko:_ wso:
rkb:1
s:b2
Done s:2F78 b1
Done s:2F9C b2
Threads finished: (4):4, (5):5, shared balance:6000
$ _
```